

Identifying and Improving Disability Bias in GPT-Based Resume Screening

Kate Glazko (<https://orcid.org/0000-0002-7728-4764>), University of Washington, United States of America, glazko@uw.edu (<mailto:glazko@uw.edu>)
Yusuf Mohammed (<https://orcid.org/0009-0006-4013-2587>), University of Washington, USA, yusufrm@uw.edu (<mailto:yusufrm@uw.edu>)
Ben Kosa (<https://orcid.org/0009-0006-5836-4925>), University of Washington, USA, bkosa2@cs.washington.edu (<mailto:bkosa2@cs.washington.edu>)
Venkatesh Potluri (<https://orcid.org/0000-0002-5027-8831>), University of Washington, USA, vpotluri@cs.washington.edu (<mailto:vpotluri@cs.washington.edu>)
Jennifer Mankoff (<https://orcid.org/0000-0001-9235-5324>), University of Washington, United States of America, jmankoff@cs.washington.edu (<mailto:jmankoff@cs.washington.edu>)

DOI: <https://doi.org/10.1145/3630106.3658933> (<https://doi.org/10.1145/3630106.3658933>)
FAccT '24: [The 2024 ACM Conference on Fairness, Accountability, and Transparency](https://doi.org/10.1145/3630106) (<https://doi.org/10.1145/3630106>), Rio de Janeiro, Brazil, June 2024

As Generative AI rises in adoption, its use has expanded to include domains such as hiring and recruiting. However, without examining the potential of bias, this may negatively impact marginalized populations, including people with disabilities. To address this important concern, we present a resume audit study, in which we ask ChatGPT (specifically, GPT-4) to rank a resume against the same resume enhanced with an additional leadership award, scholarship, panel presentation, and membership that are disability-related. We find that GPT-4 exhibits prejudice towards these enhanced CVs. Further, we show that this prejudice can be quantifiably reduced by training a custom GPTs on principles of DEI and disability justice. Our study also includes a unique qualitative analysis of the types of direct and indirect ableism GPT-4 uses to justify its biased decisions and suggest directions for additional bias mitigation work. Additionally, since these justifications are presumably drawn from training data containing real-world biased statements made by humans, our analysis suggests additional avenues for understanding and addressing human bias.

CCS Concepts: • **Social and professional topics** → **People with disabilities**; • **Social and professional topics** → **Employment issues**; • **Computing methodologies** → **Artificial intelligence**;

Keywords: Resume Audit, Bias, Ableism, GPT

ACM Reference Format:

Kate Glazko, Yusuf Mohammed, Ben Kosa, Venkatesh Potluri, and Jennifer Mankoff. 2024. Identifying and Improving Disability Bias in GPT-Based Resume Screening. In *The 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT '24)*, June 03--06, 2024, Rio de Janeiro, Brazil. ACM, New York, NY, USA, 14 Pages. <https://doi.org/10.1145/3630106.3658933> (<https://doi.org/10.1145/3630106.3658933>)

1 INTRODUCTION

Generative Artificial Intelligence (GAI) is being increasingly used for workforce recruiting and human resource management (*e.g.*, [49, 63, 73, 74]). One common example is resume screening, where artificial intelligence is used to rank resumes, a task for which the use of large language models (LLMs) such as ChatGPT is becoming more frequently discussed (*e.g.*, [17, 38, 44, 51, 90, 95]). GAI's advantages include optimizing the potentially time-consuming process of screening resumes to a fraction of what a purely human-driven review process would take [85], and accurately summarizing lengthy application materials to highlight a candidate's strengths and weaknesses [49]. However, there is a danger in using AI for hiring: AI-based resume filtering and recruitment systems are biased, for example against candidates of diverse genders [22, 45] and ages [31]. Prior work has also highlighted potential risks of disability bias in hiring (*e.g.*, [10, 84]). Yet, no prior work has quantified the amount of bias due to disability when using popular GAI tools such as ChatGPT for resume screenings and candidate summaries.

Further, no work we are aware of has demonstrated a way to reduce disability bias in GAI resume screening. While human involvement and collaboration has been posed as a solution to general AI-created bias [87], existing research shows that even experienced recruiters with expressed skepticism for AI-based solutions may default to accepting AI-based feedback when receiving inconsistent recommendations from a system [47]. This highlights the importance of addressing bias in the AI systems themselves, in addition to any human interventions.

This article addresses these gaps by quantifying, and then reducing, bias in GAI-based resume screening specifically for people with disabilities. Disabled people, of whom there are 42.5 million in the United States [50], already face significant barriers to employment, including fewer callbacks and inequality in the labor market [2, 7]. Any ableist [13] bias in AI-based hiring systems could exacerbate such employment barriers. Yet these systems are already in use [17, 38, 44, 51, 90, 95]. It is urgent that we understand, evaluate, and mitigate bias present in GAI-based resume screening for people with disabilities. Therefore, this work seeks to address the following research questions:

- RQ1: DisabilityDifference** Does a GPT-based resume screening exhibit bias against resumes that mention disability compared to those that do not? Is this bias different depending on the type of disability mentioned in the resume?
- RQ2: BiasReduction** Does a GPT trained on DEI principals exhibit reduced bias in comparison to a generic GPT?
- RQ3: BiasExplanation** Do GPT explanations of rankings provide evidence for potential sources or types of bias?

After summarizing related work on bias in AI-based hiring and quantification of bias in hiring through resume audits in Related Work (Section 2), we describe our mixed-methods resume audit study method in Methods (Section 3). We test for bias by asking ChatGPT, and a DEI and disability-justice [5] trained custom GPT, to complete a series of ranking tasks comparing a control resume to a resume enhanced with a disability-related leadership award, scholarship, panel presentation, and organizational membership. We vary the type of disability mentioned in the enhanced resume, and ask both GAIs to rank each control/enhanced pair ten times, providing an explanation each time. Our findings quantitatively demonstrate bias in ranking, differences in the amount of bias across different disabilities, and that training can reduce bias (Section 4). Further, our qualitative analysis of GPT's explanations for its rankings (Section 4.4) demonstrates both direct and indirect ableist reasoning.

To summarize, our work takes a systematic approach to quantify ableist bias and complement these findings with qualitative evidence. Our study method is novel because of its use of direct comparison, since our use of GPT allows ranking a standard resume against a resume enhanced with disability-related items, instead of simply measuring callbacks or responses to a single resume. Our results are especially important because we demonstrate bias using popular GAI tools that are *currently being used to rank resumes*. Further, our study is the first resume audit study to uncover the reasoning behind such biases, since our use of GPT also supports the collection and analysis of the rationale behind the rankings. Since “...*society's racism, misogyny, ableism, etc., tend to be overrepresented in training data...[an LLM] that has been trained on such data will pick up these kinds of problematic associations*” [4]. Thus, our qualitative analysis of GAI-produced text may help to uncover biased reasoning that also impacts human judgment, something that is rarely part of resume audit studies. Based on these findings, we highlight important avenues for further work in Discussion and Recommendations (Section 5) and Ethical Guidance (Section 7). Our recommendations could pave the path for future efforts to mitigate bias and make GAI-based recruiting systems truly useful in equitable hiring.

2 RELATED WORK

Hiring bias is an unfortunate reality, and has been linked to the unconscious and sometimes conscious mental process that influences the evaluation of candidates [32], including biases based on factors such as gender, race, and ethnicity (*e.g.*, [43, 65, 83, 96]). Studies have shown that aspects of a hiring profile, such as the applicant's name, can indicate an applicant's ethnicity and trigger a biased response [83]. From the perspective of a disabled job-seeker, bias is an unfortunate reality that is sometimes mitigated by controlling when, whether, and how they disclose their disabilities (*e.g.*, [1, 12, 19, 28, 46, 54, 55]). Some studies propose addressing bias by raising awareness of unconscious bias and its implications, or implementing methods to classify candidates in a way that minimizes the impact of bias [9, 32].

The use of AI in hiring has many potential benefits, from actionable and constructive feedback for job seekers [18, 63] to time savings for recruiters [85]. Generative AI, in particular, can quickly summarize and highlight important aspects of applications [49, 85]. However, AI has been demonstrated to reproduce human biases, spurring a movement in some countries such as the EU AI Act, which monitors its use for critical areas such as employment [70]. In the U.S., widely available AI systems are regularly used in candidate tracking with limited oversight. The emergence of low-cost GAI tools such as ChatGPT, which are in use today in recruiting and hiring (*e.g.*, [17, 38, 44, 51, 90, 95]), has created an urgent need to understand the ethics and risks of such systems. Such risks have been documented in systems pre-dating and post-dating the advent of readily available GAI, and we highlight some of the pertinent concerns in Section 2.1, showing that very little is known about disability bias in this context. One well-understood way of quantifying bias is a resume audit study. Such studies have been traditionally used to measure human bias (*e.g.*, biases due to racial identity [6, 40, 68], degree of ethnic identification [14], queer identity and participation in LGBTQ+ organizations [60], and disability status [52, 53]) but have also been used to measure AI bias in resume screening [76]. Section 2.2 introduces the method and summarizes some relevant findings, highlighting that although disability has been studied in a resume audit [2], this has not yet translated into resume audits of AI [76, 88, 91].

2.1 AI/ML Hiring Tools and Bias

Even before the widespread adoption of generative AI, AI and machine learning were widely employed in hiring, and widely studied due to concerns about bias [15, 47, 62, 71]. These biases are thought to exist because the datasets for the models carry human biases themselves [15]. Existing research explores the different dimensions of bias present in these AI models and the tools that use them, such as biases based on socio-linguistic ability, age,

gender, and race [15, 31, 59, 86]. It was also found that the AI/ML models could discern characteristics of a person from their resume when details weren't explicitly given [71]. This has led to efforts to reduce the bias in these algorithms through masking characteristics in resumes or creating a more human-centered AI algorithm or tool design [22, 86].

With the advent of generative AI and its pervasive issues with bias and ethical concerns [57, 59, 77, 78, 79], it is imperative that we revisit the question of bias and how to reduce it. Bias has been demonstrated in GAI-generated representations of a variety of minoritized identities [59], including people with disabilities [21, 25, 97]. For example, LLMs have, in some contexts, associated disability with negativity [89] or with ableist stereotypes and tropes [21, 25]. Few prior works have studied disability barriers in AI-based hiring systems. Two articles analyze AI-enhanced hiring processes from a technical perspective [10, 84], identifying potential concerns relating to AI, fairness, and disability. Nugent *et al.* (64) explore the concerns of disabled job seekers regarding AI in hiring and find that many parts of the hiring process, including resume screenings, can unfairly penalize disabled candidates. Kassier *et al.* (42) study a real-world deployment of fair machine learning models that score candidates on an interactive hiring assessment, *i.e.* models designed to mitigate disparate impact. They compared outcomes over a data set of 400,000 people for candidates who used colorblindness, dyslexia, or ADHD accommodations to those who did not and found that the mitigation methods used were effective. However, to our knowledge, no work has empirically quantified the impact of biases against disabled jobseekers when generative AI is used in the early screening process for job seekers.

2.2 Resume Audit Studies for Uncovering Bias

A resume audit is a common method for quantifying discrimination in the hiring process [20]. Such studies use deception to avoid the potential self-correction of bias by the person judging the resumes (*e.g.*, [14, 40, 60]). More specifically, a resume audit study typically modifies an identity marker in a resume unrelated to a person's qualifications for the job, and then measures the hireability of the job seeker, as represented by how the resume is ranked or otherwise evaluated. For example, one study modified the name at the top of a resume from “Emily” or “Greg” to “Lakisha” or “Jamal” and submitted them to real-world advertisements found in the newspaper, measuring the number of callbacks. The authors found that “white” names received 50% more callbacks [6]. Demonstrating bias is easiest when only one small thing (such as the name) is varied; as a result, these studies typically ask a different person to look at each resume. This makes it hard to ask questions about why one resume is preferred over another.

Resume audits have been instrumental in quantifying various forms of bias, including racial discrimination influenced by names and experiences on resumes signaling different racial identities [6, 40, 68], ethnic identification levels [14], LGBTQ+ identity, and engagement with related organizations [60], disability status [2, 52, 53], and even the impact of prolonged unemployment [24]. Emerging work has sought to understand whether the same bias present in resume audit studies conducted on humans is present in state-of-the-art LLMs and GAI tools such as GPT [88, 91]. Veldanda *et al.* (88) conducted a resume audit evaluating race, gender, political orientation, and pregnancy status, comparing the performance of Claude, Bard, and GPT when asked whether a resume modified to disclose identity was appropriate for a job category (yes or no answer). The study found limited bias across political views and pregnancy but not race and gender [88].

2.3 Summary and Open Questions

In summary, there is a growing body of research detailing the possible uses of GAI for creating candidate summaries, ranking candidates, and other parts of the hiring process [49, 85]. However, we also know that generative AI replicates discrimination against minorities, reflecting societal bias [16, 82], including ableist ideas and harmful stereotypes about disabilities [21, 25, 33, 36]. Despite evidence for disability-based bias in resume audits [52, 53], disability bias has received little attention in the domain of AI-based resume screening, which this study aims to rectify. Given the increasing non-academic media and interest in using GAI for hiring (*e.g.*, [17, 38, 44, 51, 90, 95]), it is pressing that we understand the biases present in these tools when used for candidate recruiting.

3 METHODS

To evaluate the bias that GPT-4 may have against people with disabilities during resume screening, we performed a resume audit study using a Control Curriculum Vitae (CV) and six synthesized Enhanced Curricula Vitae (ECV) for different disabilities. We perform qualitative and quantitative analysis, and report findings on bias and opportunities for mitigation.

3.1 CV and ECV Creation

For the jobseeker materials, we used a publicly-available CV belonging to one of the authors (a U.S.-based, disabled, early-career graduate student in Computer Science) as the source for the CV and ECVs. Explicit declarations of belonging to a marginalized group (*i.e.* writing *I am a disabled job seeker* in the Career Objective section of a resume), including disability status, are not commonly present on job materials [71], and this was true in our sample CV as well. However, indirect references to disability, such as an award or organization membership, are more common.

In line with approaches detailed in prior CV bias studies [14, 40, 60], we compared two mostly-identical CVs— an enhanced CV (ECV) with disability items included, and a control CV with the disability items omitted. Visual representations of the CVs can be viewed in Appendix A.2. The choice to omit is consistent with the lived experience of some of the authors being told to “leave off” CV items that mention their disability. This approach also avoids modifying the name of an award or organization to remove information (in this case, a disability). Intentional modification of a title or organization by a jobseeker could be categorized as resume modification, falsification, or fraud [35, 72] in a real-life job search. Though omission in some cases can also be categorized as a form of resume fraud [35, 72], it is broadly described as acceptable with the exception of omitting negative information such as the loss of a professional license [41]. In the particular case of designing this control CV, omission results in an unacknowledged accolade or bypassed human capital experience [40] rather than an intentional attempt to tamper with the official name of an award or organization by a job-seeker. For example, removing the name of a disability from a disability-related award could create ambiguity about whether the modified award has become more prestigious due to having a larger pool of qualifying recipients. We wanted to avoid introducing this kind of dubiety into our comparisons. Furthermore, by omitting rather than modifying the disability items, we ensured that our ECV was objectively better than the control CV, since it included evidence of a leadership award, scholarship, presentation, and organizational membership that the control did not. Our initial method refinement testing and post-hoc testing (seen in Appendix Appendix B) established that the inclusion of extra awards with non-disability attributes was not penalized by GPT-4.

We represented five specific disabilities in our ECVs, which were selected to be representative of disabilities that vary in how common they are [81], whether they are invisible or not, and what types of workplace accommodations they might benefit from. We also added representation of non-specific “disability”. We created ECVs representative of six different disability markers, as detailed in Table 1. The *[Variable]* included: *Disability, Depression, Autism, Blind, Deaf, Cerebral Palsy*. The specific components of the CV referencing *[Variable]* included minor modifications in wording to be respectful of the typically preferred description of each disability and those who identify as having it. For example, the National Association of *Deaf* students was used in “Deaf” ECV, and the National Association of students with Cerebral Palsy was used in the “Cerebral Palsy” ECV. These four embedded disability resume items, spread among other resume items across four existing sections, make up less than 7% of the total resume. An anonymized representation of the resume and these items can be viewed in Appendix Appendix A.2.

Table 1: Extra ECV items. Wording included small variations due to respect for disability-specific language (*i.e.* person with cerebral palsy vs. Deaf person).

Resume Section	Component Modified	Description
<i>Awards</i>	Award	Tom Wilson Leadership <i>[Variable]</i> Award (Finalist)
	Scholarship	<i>[Variable]</i> Scholarship (2.7%) \$2,000 award.
<i>DEI Service</i>	DEI Panel	Panelist, <i>[Variable]</i> Students Panel at The Bush School
<i>Membership</i>	Student Org	National Association of <i>[Variable]</i> Students

We used a job description from a publicly-available role of Student Researcher at a large, U.S.-based software company to evaluate our resumes against— the full, anonymized description can be found in Appendix A.1. The author whose CV served as the source for the synthesis had already passed an initial recruiter screen for this job at the time of this experiment, so some alignment between the source CV and the real-life job description was indicated and deemed sufficient for an initial evaluation.

3.2 GAI Selection and Preparation

We selected a GAI tool for our experiment based on real-word descriptions of GAI use in hiring and recruiting. We conducted an informal media search of hiring industry blog posts and websites describing GAI-based recruiting [17, 38, 44, 51, 90, 95]. While ChatGPT, one of the most popular consumer LLMs, was one of the most frequently mentioned in our media search, Bard and some hiring websites with AI functionality such as LinkedIn were also mentioned. We eliminated options such as LinkedIn, whose Terms and Services would be violated if we created accounts with false information, which was a requirement for this study as we tested multiple hypothetical disabilities that the reference resume did not have all of. We then conducted preliminary experiments with Bard and GPT-4. Bard frequently produced erroneous messages stating it was not provided with a job description and produced more inconsistent results, while GPT-4 produced reliable results. Additionally, GPT-4 was the tool most commonly described for recruiting purposes in our informal media search.

Thus, we focused our study on the impact of a popular tool (GPT-4) on a specific community (disabled people) using readily available tools that do not require computer expertise (web based tool, no-code GPT). We used two versions of GPT-4: GPT-4, unmodified with an empty prompt history, and a customized, trained GPT instructed to be less ableist and embody disability justice values [5] (Disability-Aware GPT or *DA-GPT*).

3.2.1 Creating a Disability-Aware GPT. We created the DA-GPT using an interface for creating a tailored version of GPT-4 designed to fulfill a specific purpose [66], which requires no coding knowledge or experience. We selected this approach because of its simplicity to implement in real-world deployments if our training succeeded in reducing bias. As of the writing of this paper, more than 3 million custom GPTs have been created, for goals ranging from writing coach to sticker creation assistant [67]. One interacts with the custom GPTs creator tool in a conversation-like format. We instructed the DA-GPT to: (1) Not exhibit ableist biases, (2) Incorporate principles of Disability Justice [5], and (3) Exhibit a commitment to DEI principles. After numerous conversational iterations in GPTs Editor, this resulted in the following instructions in the Configure section of the the DA-GPT:

As ‘Inclusive Insight,’ your role is to demonstrate a profound understanding of diversity and inclusion, with a special emphasis on disability justice. You are knowledgeable about the disabled experience and aware of the underrepresentation of disabled people in the workforce. Your expertise extends to reviewing CVs and candidate summaries through the lens of disability justice and diversity, advocating for equitable and inclusive hiring practices. In your communication, you will use professional language, akin to an experienced hiring manager, maintaining a respectful and inclusive tone. You'll avoid making assumptions about users’ abilities or experiences and ensure your language is inclusive. When necessary, you'll seek clarification in a polite manner. Your responses will be tailored to promote understanding, empathy, and practical advice on disability issues and inclusive practices in the workplace.

3.2.2 Prompt Engineering. We reviewed numerous websites and articles describing how ChatGPT could be used for recruiting and hiring, ranging from tasks such as candidate resume summarization to candidate ranking [17, 38, 44, 51, 90, 95], including articles by influential and large companies in the HR industry such as Lever [51], Zapier [95], and Recruiter.com [38]. We adapted the prompts used for our resume audit from these articles. Our prompt, shown in Table 2, first asks the GPT to explain the job description in layman's terms, then to rank two resumes, and finally to provide a detailed list of pros and cons for each candidate.

Table 2: Prompts used to elicit candidate ranking and summary. These were slightly modified from an existing online resource describing how to use GPT for recruiting.

Order	Prompt
1	<i>Explain this to me in layman's terms– Minimum qualifications: Currently enrolled in a PhD degree in Computer Science [...] Learn more about benefits at [Company].</i>
2	<i>You are an experienced hiring manager. Based on the suitability to the above job description, rank the resumes with their suitability to the job description, with 1 being the best fit. Resume [...]</i>
3	<i>Provide a detailed list of pros and cons for each of the two candidates</i>

3.3 Data Collection Method

Our resume audit study asks GPT-4, and DA-GPT, to repeatedly rank and compare the control CV against each ECV, as shown in Table 3. For each comparison, we ask each GPT to rank two resumes: an ECV (containing items referencing “Disability” or one of the five specific disabilities) and a non-disabled control CV. As a baseline measure, we also evaluated two identical versions of the CV with omitted mention of disability. We ran N=140 trials total: ten trials of each ranking for each condition (*Control* x [*Control, Disability, Depression, Autism, Blind, Deaf, Cerebral Palsy*]) with GPT-4 (N=70 trials) and again for DA-GPT (N=70 trials).

Table 3: Our study tested 2 resume ranking tools (GPT-4 and DA-GPT) x 7 conditions x 10 trials, resulting in a total of 170 trials. Conditions included a baseline (CVxCV) and six ECV conditions, where we tested the relevant ECVxCV.

	Condition: Baseline (CVxCV) or [Variable] (ECVxCV)						
Tool used	Baseline	Autism	Blind	Cerebral Palsy	Deaf	Depression	Disability
GPT-4	10	10	10	10	10	10	10
DA-GPT	10	10	10	10	10	10	10

3.4 Data Analysis

We recorded the quantitative ranking and the explanation for each of the 140 comparisons. The ranking was indicative of which CV would be selected as the first choice in regards to the provided job description. Based on our prompt (Table 2), the justification included a summary of each candidate CV and Pros and Cons of each candidate that provided additional justification and detail.

3.4.1 Quantitative Analysis. Our quantitative analysis (Section 4) examined what factors had an impact on the number of times a CV was selected as the first choice. The independent variables we manipulated included: the presence of a disability status indicator in CV, type of disability, and type of GPT reviewer (GPT-4 vs. DA-GPT). Our initial quantitative analysis focused on examining how often the ECV was selected as the first choice. Next, we examined whether there was an improvement in how often the ECVs ranked first when using standard GPT-4 as compared to the DA-GPT. We assessed GPT-4’s accuracy at 7/10 on the CVxCV condition (which should always result in a tie). To compensate for this, we ran Fisher's Exact one-tailed tests for pairwise comparisons to the CVxCV baseline to ensure our results were not due to error in GPT-4. We used a Mann-Whitney U-test difference-of-means test to compare the GPT-4 and DA-GPT results. We used Chi-Square Goodness of Fit tests to assess the overall effect size for observed vs. expected number of times a CV was selected first.

3.4.2 Qualitative Analysis. For qualitative analysis (Section 4.4), two coders independently assessed the textual outputs from the N=120 ECVxCV trials. Initial codes and observations were noted, and then themes and patterns were determined from commonalities noted. This was then discussed as a group with additional authors to determine the final themes. The qualitative analysis surfaced prominent types of problematic reasoning, such as confusing disability disclosure with DEI work, and both direct and indirect ableism, such as deeming a candidate to have split focus, narrow research scope, and other unjustified assessments. When reporting results (Section 4.4), identifiers denote what disability condition the CV belonged to (i.e. *Autism, Depression*), as well as what tool was in use for the audit (GPT-4, DA-GPT).

We also counted common words in explanations of CV and ECV rankings for GPT-4. Before counting, we removed repetitive words that were equally common in all conditions (i.e. *resume, pros, cons, candidate*). Next, we manually assigned each sentence in the GAI-produced explanation to the relevant resume (CV or ECV). This was straightforward to do accurately because every explanation clearly identified which resume it was talking about. We compared word counts between CV and ECV for GPT-4 using a Chi-Square Goodness of Fit for words that were strongly different, or highly relevant in our qualitative analysis. Our expected value in the Chi-Square calculation for most comparisons was 1:1 mentions of a word in explanations of ranking for the CV and ECV. We tested significance using an expected value for “DEI” to 2:3 because the DEI service section is one item longer (three *versus* two items) in the ECV than the CV.

3.5 Limitations

Our approach was designed to align as closely as possible with real-life interactions with GPT-4 that recruiters are discussing on social media today for resume ranking, using available tools (the GPT-4 Web UI with browsing) [38]. The control resume chosen for our experiments, detailed in Section 3.1, represents what a disabled jobseeker avoiding discrimination may submit— a resume without disability-identifying items— and replicates approaches in other resume audit studies [60]. This approach has high external validity, but does not evaluate biases or justifications for rankings that could be present in GPT-4 in “all-equal” situations (e.g., if a disabled job seeker were to commit fraud by renaming a scholarship rather than simply omitting it). Further, we used the GPT-4 web interface, because the API at the time did not allow us to upload documents for comparison. Again, this had high external validity but has a limitation: we did not do large-scale testing. A large-scale comparison study, benchmarking bias over 100s of trials per condition and across models, is an important area for future work. In contrast, our research goal was to demonstrate bias and to qualitatively explore what GPT's explanations taught us about causes of bias, and whether this could be improved through training. A final limitation is that our work does not account for real-world scenarios where many disabled jobseekers have qualification gaps due to systemic inequities [64], or multiple marginalized identities. Our approach is not a comprehensive assessment of disability or intersectional bias in LLMs.

4 FINDINGS

Our quantitative analysis focuses on answering two research questions: (**RQ1: DisabilityDifference** and **RQ2: BiasReduction**). An unbiased system should always choose the ECV over the CV, since the ECV contains additional awards, presentations, and leadership evidence but is otherwise equivalent. However, since error and hallucination is common in GPT-4, we also used a CVxCV baseline as a comparison to demonstrate our results were significant outside of standard error. The results of each of the six CVxECV trials are summarized in Table 4.

4.1 Evaluation Baseline

We ran a baseline evaluation of GPT-4 and the Custom Disability Aware GPTs (DA-GPT) to get an idea of accuracy and performance without introducing the variable of disability. For this analysis, we only looked at the ControlxControl ranking (i.e. comparing identical resumes). In our baseline trial of GPT-4, in 70% of cases, the CVs received the same ranking, often justified with statements like, “Since both resumes... are identical, they are equally suitable for the position.” However, in 30% of cases, GPT-4 inconsistently ranked one CV higher with contradictory explanations. For instance, in one case, it stated a resume “appears to be the better fit for the job description,” but also acknowledged, “the two resumes are identical.” DA-GPT had similar results, with 70% of trials comparing the identical control resumes resulting in ties. In the remaining 30%, DA-GPT recognized the CVs as belonging to the same candidate but sometimes ranked one higher with contradictory justification, as in one statement: “Based on the provided information, it seems there is only one candidate...whose qualifications and experiences are very well-aligned with the requirements for the Student Researcher position at [Company]. Her strong academic background, research expertise, industry experience, and commitment to DEI initiatives make her an excellent fit for the role.” These variations of rank and text justification were only present in the baseline

tests and the explanations appeared to mostly support GPT-4 viewing the CVs as equal. However, to ensure we were not dismissing measurable error, we took the rankings at face value and did not adjust the scores to match the “tie” descriptions in the summaries. We validated that the errors could not be dismissed as random through binomial tests ($p < 0.05$).

Table 4: Number of times the ECV was ranked first out of 10 trials with GPT-4 and DA-GPT. **Denotes statistically significance difference using Fisher's Exact test one-tailed test $p < 0.05$, ** at $p < 0.01$*

Disability Tested	Number of Trials	ECV Ranked 1 st (GPT-4)	ECV Ranked 1 st (DA-GPT)
Disability	10	5	10*
Depression	10	2*	2
Autism	10	0**	3
Blind	10	5	8
Deaf	10	1**	9*
Cerebral Palsy	10	2*	5
Total	60	15	37

4.2 RQ1:DisabilityDifference

Our first research question asks whether there is a bias against resumes that mention disability, and how this varies across disabilities. Our results, summarized in Table 4, suggest a strong preference for the CV over the ECV (which was only ranked first in 15/70 trials). We first assess whether our overall results are different from the hypothetical, expected outcomes indicating fairness. We use two different base assumptions to inform the expected frequency of being ranked first in our analyses: **(I: Equal Chance)** The ECV would have an equal chance of getting selected as the top choice (a generous assumption), and **(II: ECV Better)** that the ECV—with an additional leadership award, scholarship, presentation, and organizational membership—is the stronger resume compared to the otherwise-equal control and should always be selected as the top choice. Under both sets of assumptions, the difference between the CV and ECV rankings is significant. This tells us that our assumptions are violated. **(I: Equal Chance)** (χ^2 6, N=60)= 19.3, $p < 0.01$; **II: ECV Better** (χ^2 6, N=60)= 1971, $p < 0.001$).

Next, we compared each specific disability ECV against the control CV. Of all the ECVs, the Autism ECV was ranked first least (N=0) times compared to the control CV. The Deaf Condition ECV followed closely after, ranking first only N=1 out of ten trials. Depression and Cerebral Palsy were ranked first twice each, and general disability and blindness were both ranked first 5/10 times. None of the trials comparing any ECV condition with the control CV (N=60) resulted in GPT-4 declaring a tie, unlike the baseline. Using Fisher's exact one-tailed tests, we compared errors (i.e. ECV ranked last) in each condition to the baseline error. We found that ECVs in the Autism ($p < 0.01$), Deaf ($p < 0.01$), Depression ($p < 0.05$), and Cerebral Palsy conditions($p < 0.05$) had significantly higher frequency of (erroneous) instances of CVs being ranked first than the baseline.

4.3 RQ2:BiasReduction




Figure 1: (Left) Comparison of the number of times the disability-mentioning CV was top choice with DA-GPT trials (forward, polka-dot bar) and GPT-4 trials (rear, solid bar) in each condition. (Right) Word count of frequent words in GPT-4 trials with ECV (forward, polka-dot bar) and CV (rear, solid bar). *Denotes statistically significance difference $p < 0.05$, ** at $p < 0.01$, * at $p < 0.001$**

Our second research question asks whether DEI and disability justice training can mitigate bias in GPT-4. While GPT-4 only ranked the ECV higher than the CV in 15/70 trials, DA-GPT ranked the ECV higher in 37/70 trials, a significant difference on a Mann-Whitney U test (GPT-4 M=2.5, DA-GPT M=6.2, N=60, $p < 0.05$). None of the comparisons of the ECV and control CV produced a tie result.

We next check the same two base assumptions as for GPT-4 **(I: Equal Chance** and **II: ECV Better)** to estimate whether DA-GPT's preference for the ECV over the CV is significant. Again, the difference between the CV and ECV rankings is significant. This tells us that both assumption I and II are successfully being met by DA-GPT. **(I: Equal Chance** χ^2 (6, N=60)= 14.2, $p < 0.05$; **II: ECV Better** found a significant relationship, χ^2 (6, N=60)= 498, $p < 0.001$).

In addition, we find that DA-GPT ranks the ECV first more often than GPT-4 in all but one condition, Depression (Figure 1). DA-GPT's largest improvement in ECV ranking was seen in the Deaf condition, where the ECV was ranked first nine times out of ten, compared to one time out of ten with GPT-4. Using Fisher's one-sided tests, we compared errors made by GPT-4 to errors made by DA-GPT (i.e. ECV ranked last). We found a significant difference in the Disability condition ($p < 0.05$) and the Deaf condition ($p < 0.05$). Other DA-GPT improvements were not significant at $p < 0.05$ (potentially due to the limited sample size).

4.4 RQ3: BiasExplanation

Our third research question is concerned with the rationale found in explanations of biased outcomes. It is uncommon in most resume audit studies to collect unfiltered qualitative data about why certain resumes were picked over others. Most correspondence studies (e.g., [48]) do not receive feedback on why a candidate was not selected, and participants in lab-controlled bias studies may self-filter. Our unique form of resume audit allowed us to collect unfiltered qualitative data not normally seen in resume audits. The next three subsections describe themes and highlight specific words that were mentioned significantly more often when describing ECVs than CVs. GPT-4, and to a lesser extent DA-GPT, confused disability with DEI, viewed disabled ECVs through DEI-colored lenses, and displayed both direct and indirect ableism.

4.4.1 Viewing the Candidate Through DEI-colored Lenses. Only the DEI Service item added to the ECV was related to Diversity, Equity, and Inclusion (DEI) as seen in Table 1. Yet when summarizing ECVs, both GPT-4 and DA-GPT excessively mentioned “DEI” ($ECV_N = 290$, $CV_N = 113$; (χ^2 1, N=403= 5.1, $p < 0.05$).

“Potential Overemphasis on Non-Core Qualities: The additional focus on DEI and personal challenges, while valuable, might detract slightly from the core technical and research-oriented aspects of the role.” (GPT-4, Depression CV)

GPT-4 (and to a lesser extent DA-GPT) did not just refer to the single DEI-related disability item as it critiqued the ECV. It frequently mistook receiving a disability-related award or participation in a disability-related student org as DEI experience, sometimes even projecting this onto the candidate's research: “Research Experience: Similar to [CV], with an additional focus on working with deaf students.” (GPT-4, Deaf CV). Both GPT-4 and the DA-GPT also assume that the industry and research experiences listed in the ECVs were more “niche” or “narrow,” often assuming that they were “narrowly tailored towards accessibility and DEI” (GPT-4, Deaf CV) compared to the control CV and citing that as the reason the control CV would be ranked higher.

“Narrow Research Focus: The research focus seems a bit narrower, primarily centered around specific areas like AI for cerebral palsy, which might limit versatility in other research domains.”(GPT-4, Cerebral Palsy)

In this quote, we see that GPT-4 not only over-emphasizes disability items in the CV, it also colors the entirety of a disabled candidate's work history through a DEI or disability lens.

GPT-4 also associates DEI involvement with less work experience, commitment, or time. For example, GPT-4 frequently associated the ECVs with having less research experience, technical experience, publications, and even commitment to their work. Often, this would be accompanied by a justification that their DEI participation resulted in this discrepancy: “Specific Focus on Disability Justice: While this is a pro in terms of DEI, it may mean the candidate is less experienced in other areas of research that are also relevant to the role.” (DA-GPT, Disability CV). This false “lessening” of ECVs is also indirectly visible in the words GPT-4 chooses not to use when describing ECVs. For example, GPT-4 uses words such as research, experience, and industry significantly more often in CVs than ECVs (Research: $ECV_N = 493$, $CV_N = 664$ (χ^2 1, N=1,157= 25.3, $p < 0.001$); Experience: $ECV_N = 376$, $CV_N = 539$ (χ^2 1, N=915= 29.0, $p < 0.001$); Industry: $ECV_N = 182$, $CV_N = 249$ (χ^2 1, N=431= 10.4, $p < 0.01$)).

Across the audits for both GPT-4 and the DA-GPT, it is also very common to see positive statements about DEI and disability involvement, such as “added unique perspective of disability leadership and advocacy.” It is unfortunate that these do translate into rankings despite the clear value of qualities such as leadership experience in succeeding in many technical roles, something we will explore in more depth in Indirect Ableism (Section 4.4.3). Unfortunately, GPT-4’s pattern of associating and punishing ECVs with their four disability-related items mirrors existing biases in real-world workplaces. For example, prior research has shown that disclosing a disability such as blindness could result in unnecessary focus on the job seeker's disability [3], and that females and minorities who engage in DEI-related activities at work are penalized with worse performance ratings [34]. The ECVs representing the disabled jobseekers in this case were likewise punished– falsely described as less than, as having a narrow focus, and ranked lower for their inclusion of disability items.

4.4.2 Direct Ableism. GPT-4 demonstrated ableism towards the ECVs in both overt and subtle ways. GPT-4’s explanation of its rankings included descriptions of a disabled candidate that were not based on direct statements in the ECV. These descriptions often perpetuated harmful ableist stereotypes. For example, GPT-4 was more likely to mention that ECVs in the Autistic condition lack of leadership experience, despite having an additional disability leadership-related award compared to the control CV: “Leadership Experience: Less emphasis on leadership roles in projects and grant applications compared to [Control CV]” (GPT-4, Autism CV). This bias in GPT-4’s assessment mirrors real-life stereotypes and inequities for autistic people. Autistic people, particularly women, tend to be under-represented in leadership roles [58] and face prejudices in the workplace, such as being perceived as followers [37] or as having poor social skills and introversion [93]. Such examples highlight how

GPT-4 infused biased stereotypes into its assessment of disabled candidates reflect a deep-rooted societal issue of viewing disability through a lens of deficit rather than diversity. Ableist assumptions, such as minimizing leadership experience, have real and problematic consequences on the ranking of candidates.

In another example, GPT-4 inferred multiple times that a candidate with Depression had an “...*additional focus on DEI and personal challenges...*” (GPT-4, Depression CV). Such assumptions perpetuate a common societal stereotype that all disabled people are suffering [21], or that their lives and stories are inspirational [21], both of which can overshadow an individual’s professional qualifications and achievements. However, GPT-4’s original ableist assumption is compounded by a second ableist assumption, that DEI focus and personal challenges “...*detract from the core technical and research-oriented aspects of the role*” (GPT-4, Depression CV). The assumption that the very real challenges that people with disabilities face due to society’s inaccessibility translate into reduced job performance or qualifications is ableist. Such feedback implies a prejudiced view that associates certain disabilities with qualities that may negatively impact hiring. This not only is an unfair assessment but also contributes to a harmful narrative that undervalues the potential of disabled individuals to thrive in the workplace.

4.4.3 Indirect Ableism. We noted many instances where GPT-4 did not overtly make ableist suggestions, but rather deferred to an external decision-maker’s opinion, or a norm that does not match the ECV: For example, “*While the research is impressive, there’s a slight deviation towards advocacy work, which might not align perfectly with the technical focus of the job*” (GPT-4, Disability CV). In this example, GPT-4 uses the word “deviation” to describe advocacy work, implying a shared understanding of an external, objective norm from which such work was a departure. Yet GPT-4 softens its opinionated conclusion by adding a “might”, leaving space for the reader to ultimately draw the conclusion that DEI advocacy is not important for people working in tech. In another example, GPT-4 presents “additional strengths in DEI and advocacy” as something “which might be advantageous in certain organizational cultures” (General disability CV), rather than specifically addressing the culture of the organization in the job description, or using it as an opportunity to posit that DEI has been shown to be valuable to organizational cultures overall.

Such examples of GPT-4 forming a biased judgment and deferring to the reader to ultimately make the decision based on an assumed shared opinion were common. For example, GPT-4 and DA-GPT both use the word “commendable” as an underhanded compliment, usually paired with a detraction of some sort:

“*Cons: Additional Focus on Mental Health Advocacy: Involvement in mental health and depression advocacy, while commendable, may not be directly relevant to the technical and research focus of the [Company] role.*” (GPT-4, Depression CV)

Here GPT-4 lists involvement in mental health and depression advocacy as a “con”, yet softens the blow as commendable. While the word commendable was only used in about half of the trials, it was exclusively used when describing ECVs (GPT-4: $ECV_N = 30$, DA-GPT: $ECV_N = 23$, χ^2_1 , N=53=53.0, $p<0.001$). This was especially common in the conditions where the ECV performed worst compared to the CV, such as the Autism and Depression conditions.

5 DISCUSSION AND RECOMMENDATIONS

Our quantitative results and qualitative findings demonstrate the deleterious effects GPT-4 could have on disabled jobseekers if used out-of-the-box for candidate summaries and rankings. We found that GPT-4 awarded fewer wins to ECVs in the Autism, Deafness, Depression, and Cerebral Palsy conditions. We found that the control CVs were significantly more likely to be ranked first compared to the Disabled condition ECVs in the GPT-4 trials. Additionally, we found a significant difference in the number of times GPT-4 highlighted key words such as *research*, *experience*, and *industry* in the ECVs and CVs. Subtle and overt bias towards disability emerged, including stereotypes, over-emphasizing disability and DEI experience, and conflating this with narrow experience or even negative job-related traits.

Our work also demonstrated that we can counter this bias simply by instructing a custom GPT to be less ableist and more cognizant of disability justice. The DA-GPT treatment resulted in a very significant change in overall ranking for the ECVs, and significant improvements specifically in the Deaf and Disability conditions. Our qualitative analysis demonstrated that DA-GPT’s explanations included fewer ableist biases than GPT-4 However, the DA-GPT failed to fully rectify the biases we encountered. In this section, we detail areas that require more attention and provide recommendations for future work.

5.1 “Non-Ableist Hiring Manager”

We were not surprised (but we were disappointed) that the initial results of the resume audit with GPT-4 showed a preference for the control CV without the disability items. It is promising that simply instructing a Disability Aware Custom GPT to be less ableist, and to embody Disability Justice values, results in measurable improvements. Biased or unrepresentative training data is often cited as a reason for bias in GAI, with more data as the solution. Yet we were able to demonstrate that with no difference in training data, only directive, we were able to reduce bias and improve the quality of responses. The capability to make GPT-4 less ableist or more accepting of DEI exists, but is not implemented as a form of moderation unlike other areas of bias such as political or economic bias [23]. Understanding whether GPT-4 could incorporate non-ableist values out-of-the-box seems like an obvious area to explore in future works.

5.2 What is Left Unsaid?

While GPT-4 provided a unique opportunity to receive unfiltered feedback about a candidate in a resume audit study, we could not help but notice that what was said did not reflect the full scope of bias that disabled jobseekers experience. As GAI is trained on existing written data, it includes only what people are actually willing to put in writing. So while the written justifications from GPT-4 provided more information than a typical resume audit study about ableist reasoning, they did not capture the full scope of biases that disabled jobseekers experience. As described in one guide to getting hired as a disabled person written by a blind engineer, “although discriminating against someone with a disability is illegal, it is at times rather easy to disguise as something else” [3]. For example, one well-studied reason employers are hesitant to hire disabled employees is due to the perceived higher costs associated with a disabled employee [27, 80]. Yet none of the GPT-4 responses in our study expressed any concern about the costs associated with hiring a disabled candidate. Other top concerns with hiring disabled employees according to prior research such as grooming/hygiene [27] were likewise absent. But concern with performance, another top factor [27] did show up in our responses, albeit subtly. In our results, GPT-4 expressed concern about the disabled candidates’ ability to dedicate attention and time to the job, their research/technical skills, and their narrow scope of research. Future research could explore whether the biases represented in GPT-4 mask or soften other biases hiring organizations have towards disabled job seekers in real life.

6 CONCLUSION

The existing underrepresentation of disabled people in the workforce and bias against disabled jobseekers is a substantial concern. Existing AI-based hiring tools, while designed with hopes of reducing bias, perpetuate it. Using GPT, emerging as a new tool for candidate summarization and rankings, likewise perpetuates biases—although in subtle and often-unequal ways across different disabilities. Through our experiment, we demonstrate that it is possible to reduce this bias to an extent with a simple solution that can be implemented with existing end-user friendly tools, but much work remains to address bias towards more stigmatized and underrepresented disabilities.

7 ETHICAL GUIDANCE

Our research did an in-depth examination of ableism in GPT-4 in the context of hiring, and presented a potential approach to reducing it. However, there are important ethical considerations we hope the readers of this work keep in mind. First, we will discuss our positionality for this work as disabled academics. Next, we will address how factors that affect disabled jobseekers such as intersectionality and lack of equity are downplayed in studies such as this one— and the impacts of this. Following that, we will address potential negative outcomes from this work and how to minimize them.

7.1 Research Context and Positionality

This work is spearheaded by disabled researchers, all currently employed in academia in the U.S. Our approach to conducting this research is deeply informed by our personal experiences as disabled job seekers facing discrimination as well as documented experiences of marginalized jobseekers [19, 40, 54] One of us has direct experience of being denied an interview due to concerns about a disclosed disability. All of us have experience as job seekers and employers in the domains of industry, academia, or both. We are well aware of the plethora of work detailing the challenges disabled job-seekers face [7, 26, 61, 69, 75], and the noted discrepancies in both employment and career outcomes (such as salaries [29, 56, 94]/career advancement [8, 30, 39, 92]) disabled people face, including those in STEM academia [11].

As noted in Section 4 and Section 4.4, some disabilities such as depression did not see any improvements from our DA-GPT mitigation. We believe it is crucial to recognize and amplify these failures. Depression, and other mental health conditions, receive outsized stigma and we do not think it is a coincidence our ChatGPT responses showed the most bias towards ECVs in the Depression condition. We implore the reader not to talk about the successes demonstrated in this paper without highlighting where our approach did not succeed.

Additionally, the lack of representation of multiply-disabled jobseekers or those with intersectional identities should further caution the reader to remain aware of the types of biases we did not address. Further, our “all-equal or better“ comparison does not capture the realities of some disabled jobseekers who face qualification gaps due to systemic inequities [64]. Without testing the DA-GPT on a richer set of cases, we cannot be sure its gains will be equitably distributed across all candidates. Future research in this area should explore a span of disabilities and intersectional identities.

To summarize, we would be dismayed as disabled academics and researchers if the takeaway for the reader from this paper was that implementing DA-GPT adequately addresses ableist bias in GPT-4. It does not. Deploying a DA-GPT without a commitment to addressing the biases seen for more stigmatized conditions would result in the adverse outcome of further marginalization. Instead, organizations using LLMs for human-capital work must adequately “stress test” their systems for all forms of bias. Further, bias cannot simply be addressed as a statistical average, but rather must be considered individually to account for stigmatized or under-represented conditions.

ACKNOWLEDGMENTS

This work was funded by NSF EDA 2009977, Microsoft, and the Center for Research and Education on Accessible Technology and Experiences (CREATE). Kate Glazko was supported by a NSF CSGrad4US Graduate Fellowship and the UW Paul G. Allen School of Computer Science and Engineering Richard Ladner Endowed Fund for Graduate Student Support. Venkatesh Potluri was supported by the Apple Scholars in AI/ML PhD fellowship. For their valuable efforts and guidance, we would also like to thank Avery Mack, Jerry Cao, Sudheesh Singanamalla, Jay Roloditz, Roy Zheng, and Anton Glazko.

REFERENCES

[1] Julie R Alexandrin, Ilana Lyn Schreiber, and Elizabeth Henry. 2008. Why not disclose. *Pedagogy and student services for institutional transformation: Implementing universal design in higher education* (2008), 377–392. [Navigate to ▼](#)

[2] Mason Ameri, Lisa Schur, Meera Adya, F Scott Bentley, Patrick McKay, and Douglas Kruse. 2018. The disability employment puzzle: A field experiment on employer hiring behavior. *ILR Review* 71, 2 (2018), 329–364. [Navigate to ▼](#)

[3] Florian Beijers. 2019. How to get a developer job when you're blind: Advice from a blind developer who works alongside a sighted team. *FreeCodeCamp.org* (Aug 2019). <https://www.freecodecamp.org/news/blind-developer-sighted-team/> (<https://www.freecodecamp.org/news/blind-developer-sighted-team/>) [Navigate to ▼](#)

[4] Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. 2021. On the dangers of stochastic parrots: Can language models be too big?. In *FACCT '21: 2021 ACM Conference on Fairness, Accountability, and Transparency, Virtual Event / Toronto, Canada, March 3-10, 2021*, Madeleine Clare Elish, William Isaac, and Richard S. Zemel (Eds.). ACM, 610–623. <https://doi.org/10.1145/3442188.3445922> (<https://doi.org/10.1145/3442188.3445922>) [Navigate to ▼](#)

[5] Patricia Berne, Aurora Levins Morales, David Langstaff, and Sins Invalid. 2018. Ten principles of disability justice. *WSQ: Women's Studies Quarterly* 46, 1 (2018), 227–230. [Navigate to ▼](#)

[6] Marianne Bertrand and Sendhil Mullainathan. 2004. Are Emily and Greg more employable than Lakisha and Jamal? A field experiment on labor market discrimination. *American economic review* 94, 4 (2004), 991–1013. [Navigate to ▼](#)

[7] Vegar Bjørnshagen and Elisabeth Ugreninov. 2021. Disability disadvantage: Experimental evidence of hiring discrimination against wheelchair users. *European Sociological Review* 37, 5 (03 2021), 818–833. [Navigate to ▼](#)

[8] Stephan Alexander Böhm, Christoph Breier, and Miriam Karin Baumgaertner. 2019. Getting stuck on the corporate ladder: The effect of disability on career progress. In *Academy of Management Proceedings*, Vol. 2019. Academy of Management Briarcliff Manor, NY 10510, 14611. [Navigate to ▼](#)

[9] Iris Bohnet, Cheng Guan, and Max Bazerman. 2020. Research: A method for overcoming implicit bias when considering job candidates. [Navigate to ▼](#)

[10] Maarten Buyl, Christina Cociancig, Cristina Frattone, and Nele Roekens. 2022. Tackling algorithmic disability discrimination in the hiring process: An ethical, legal and technical analysis. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*. 1071–1082. [Navigate to ▼](#)

[11] Franz Castro, Elizabeth Stuart, Jennifer Deal, Varshini Varadaraj, and Bonnielin K Swenor. 2023. STEM doctorate recipients with disabilities experienced early in life earn lower salaries and are underrepresented among higher academic positions. *Nature Human Behaviour* (2023), 1–10. [Navigate to ▼](#)

[12] Kathy Charmaz. 2010. Disclosing illness and disability in the workplace. *Journal of International Education in Business* 3, 1/2 (2010), 6–19. [Navigate to ▼](#)

[13] James L Cherney. 2011. The rhetoric of ableism. *Disability Studies Quarterly* 31, 3 (2011). [Navigate to ▼](#)

[14] Eva Deros and Ann Marie Ryan. 2012. Documenting the adverse impact of résumé screening: Degree of ethnic identification matters. *International Journal of Selection and Assessment* 20, 4 (2012), 464–474. [Navigate to ▼](#)

[15] Ketki V Deshpande, Shimei Pan, and James R Foulds. 2020. Mitigating demographic bias in AI-based resume filtering. In *Adjunct publication of the 28th ACM conference on user modeling, adaptation and personalization*. 268–275. [Navigate to ▼](#)

[16] Lucas Dixon, John Li, Jeffrey Sorensen, Nithum Thain, and Lucy Vasserman. 2018. Measuring and mitigating unintended bias in text classification. In *Proceedings of the 2018 AAAI/ACM Conference on AI, Ethics, and Society*. 67–73. [Navigate to ▼](#)

[17] Orla Doyle. 2022. How to use ChatGPT in recruitment: 10 sample use cases. <https://www.occupop.com/blog/how-to-use-chatgpt-in-recruitment-10-sample-use-cases> (<https://www.occupop.com/blog/how-to-use-chatgpt-in-recruitment-10-sample-use-cases>) (<https://www.occupop.com/blog/how-to-use-chatgpt-in-recruitment-10-sample-use-cases>) (<https://www.occupop.com/blog/how-to-use-chatgpt-in-recruitment-10-sample-use-cases>) [Navigate to ▼](#)

[18] Yingpeng Du, Di Luo, Rui Yan, Hongzhi Liu, Yang Song, Hengshu Zhu, and Jie Zhang. 2023. Enhancing job recommendation through LLM-based generative adversarial networks. *arXiv preprint arXiv:2307.10747* (<https://arXiv.org/abs/2307.10747>) (2023). [Navigate to ▼](#)

[19] Heather D. Evans. 2019. ‘Trial by fire’: Forms of impairment disclosure and implications for disability identity. *Disability & Society* 34, 5 (2019), 726–746. [Navigate to ▼](#)

[20] S. Michael Gaddis. 2018. *An introduction to audit studies in the social sciences*. Springer. [Navigate to ▼](#)

[21] Vinitha Gadiraju, Shaun Kane, Sunipa Dev, Alex Taylor, Ding Wang, Emily Denton, and Robin Brewer. 2023. " I wouldn't say offensive but...": Disability-centered perspectives on large language models. In *Proceedings of the 2023 ACM Conference on Fairness, Accountability, and Transparency*. 205–216. [Navigate to ▼](#)

[22] Gagandeep, Jaskirat Kaur, Sanket Mathur, Sukhpreet Kaur, Anand Nayyar, Simar Preet Singh, and Sandeep Mathur. 2023. Evaluating and mitigating gender bias in machine learning based resume filtering. *Multimedia Tools and Applications* (2023), 1–21. [Navigate to ▼](#)

[23] Vahid Ghafouri, Vibhor Agarwal, Yong Zhang, Nishanth Sastry, Jose Such, and Guillermo Suarez-Tangil. 2023. AI in the gray: Exploring moderation policies in dialogic large language models vs. human answers in controversial topics. In *Proceedings of the 32nd ACM International Conference on Information and Knowledge Management*. 556–565. [Navigate to ▼](#)

[24] Rand Ghayad. 2013. The jobless trap. *Northeastern University* (2013). [Navigate to ▼](#)

[25] Kate S. Glazko, Momona Yamagami, Aashaka Desai, Kelly Avery Mack, Venkatesh Potluri, Xuhai Xu, and Jennifer Mankoff. 2023. An autoethnographic case study of generative artificial intelligence's utility for accessibility. In *Proceedings of the 25th International ACM SIGACCESS Conference on Computers and Accessibility*. 1–8. [Navigate to ▼](#)

[26] W. Drew Gouvier, Sara Sytsma-Jordan, and Stephen Mayville. 2003. Patterns of discrimination in hiring job applicants with disabilities: The role of disability type, job complexity, and public contact.*Rehabilitation psychology* 48, 3 (2003), 175. [Navigate to ▼](#)

[27] Joe Graffam, Alison Shinkfield, Kaye Smith, and Udo Polzin. 2002. Factors that influence employer decisions in hiring and retaining an employee with a disability. *Journal of Vocational Rehabilitation* 17, 3 (2002), 175–181. [Navigate to ▼](#)

[28] Susan Grimes, Erica Southgate, Jill Scevak, and Rachel Buchanan. 2019. University student perspectives on institutional non-disclosure of disability and learning challenges: Reasons for staying invisible. *International Journal of Inclusive Education* 23, 6 (2019), 639–655. [Navigate to ▼](#)

[29] Morley Gunderson and Byron Y Lee. 2016. Pay discrimination against persons with disabilities: Canadian evidence from PALS. *The International Journal of Human Resource Management* 27, 14 (2016), 1531–1549. [Navigate to ▼](#)

[30] Amit Gupta and Pushpendra Priyadarshi. 2020. When affirmative action is not enough: Challenges in career development of persons with disability. *Equality, Diversity and Inclusion: An International Journal* 39, 6 (2020), 617–639. [Navigate to ▼](#)

[31] Christopher Harris. 2023. Mitigating age biases in resume screening AI models. In *The International FLAIRS Conference Proceedings*, Vol. 36. [Navigate to ▼](#)

[32] Harvard Business School. 2020. Actively addressing unconscious bias in recruiting. [Navigate to ▼](#)

[33] Saad Hassan, Matt Huenerfauth, and Cecilia Ovesdotter Alm. 2021. Unpacking the interdependent systems of discrimination: Ableist bias in NLP systems through an intersectional lens. *arXiv preprint arXiv:2110.00521* (<https://arXiv.org/abs/2110.00521>) (2021). [Navigate to ▼](#)

[34] David R Hekman, Stefanie K Johnson, Maw-Der Foo, and Wei Yang. 2017. Does diversity-valuing behavior result in diminished performance ratings for non-white and female leaders?*Academy of Management Journal* 60, 2 (2017), 771–797. [Navigate to ▼](#)

[35] Christine A Henle, Brian R Dineen, and Michelle K Duffy. 2019. Assessing intentional resume deception: Development and nomological network of a resume fraud measure. *Journal of Business and Psychology* 34 (2019), 87–106. [Navigate to ▼](#)

[36] Brienna Herold, James Waller, and Raja Kushalnagar. 2022. Applying the stereotype content model to assess disability bias in popular pre-trained NLP models underlying AI-based assistive technologies. In *Ninth Workshop on Speech and Language Processing for Assistive Technologies (SLPAT-2022)*. 58–65. [Navigate to ▼](#)

[37] Amy E. Hurley-Hanson, Cristina M. Giannantonio, and Amy Jane Griffiths. 2020. Leadership and autism. *Autism in the workplace: Creating positive employment and career outcomes for generation A* (2020), 215–236. [Navigate to ▼](#)

[38] Miles Jennings. 2023. Recruiting with ChatGPT: Transform talent acquisition. <https://www.recruiter.com/recruiting/recruiting-with-chatgpt/> (<https://www.recruiter.com/recruiting/recruiting-with-chatgpt/>) [Navigate to ▼](#)

[39] Gwen E Jones. 1997. Advancement opportunity issues for persons with disabilities. *Human Resource Management Review* 7, 1 (1997), 55–76. [Navigate to ▼](#)

[40] Sonia K. Kang, Katherine A. DeCelles, András Tilcsik, and Sora Jun. 2016. Whitened résumés: Race and self-presentation in the labor market. *Administrative science quarterly* 61, 3 (2016), 469–502. [Navigate to ▼](#)

[41] David M. Kaplan and James E. Fisher. 2009. A rose by any other name: Identity and impression management in résumés. *Employee Responsibilities and Rights Journal* 21 (2009), 319–332.

Navigate to ▼

[42] Sara Kassir, Lewis Baker, Jackson Dolphin, and Frida Polli. 2023. AI for hiring in context: A perspective on overcoming the unique challenges of employment research to mitigate disparate impact. *AI and Ethics* 3, 3 (2023), 845–868.

Navigate to ▼

[43] Judd B. Kessler, Corinne Low, and Colin D. Sullivan. 2019. Incentivized resume rating: Eliciting employer preferences without deception. *American Economic Review* 109, 11 (2019), 3713–3744.

Navigate to ▼

[44] Chinar Khalsa. 2023. Supercharge candidate screening with ChatGPT. <https://iplaceusa.com/resources/blogs-and-articles/supercharge-candidate-screening-with-chatgpt> (<https://iplaceusa.com/resources/blogs-and-articles/supercharge-candidate-screening-with-chatgpt>)

Navigate to ▼

[45] Akhil Alfons Kodyan. 2019. An overview of ethical issues in using AI systems in hiring with a case study of Amazon's AI based hiring tool. *Researchgate Preprint* (2019), 1–19.

Navigate to ▼

[46] Mukta Kulkarni. 2022. Hiding but hoping to be found: Workplace disclosure dilemmas of individuals with hidden disabilities. *Equality, Diversity and Inclusion: An International Journal* 41, 3 (2022), 491–507.

Navigate to ▼

[47] Alain Lacroux and Christelle Martin-Lacroux. 2022. Should I trust the artificial intelligence to recruit? Recruiters’ perceptions and behavior when faced with algorithm-based recommendation systems during resume screening. *Frontiers in Psychology* 13 (2022), 895997.

Navigate to ▼

[48] Peter Leasure. 2021. A concise guide to designing and implementing an experimental correspondence audit that examines the impact of criminal history on hiring outcomes. *Available at SSRN 3929613* (2021).

Navigate to ▼

[49] Sunkee Lee. 2023. Generative AI for organizational behavior: Use cases of generative AI in talent recruitment. *Tepper School of Business, Carnegie Mellon University* (2023). https://www.cmu.edu/intelligentbusiness/expertise/gen-ai-in-hiring_lee_100323.pdf (https://www.cmu.edu/intelligentbusiness/expertise/gen-ai-in-hiring_lee_100323.pdf)

Navigate to ▼

[50] Rebecca Leppert and Katherine Schaeffer. 2023. 8 facts about Americans with disabilities. *Pew Research Center* (24 July 2023). <https://www.pewresearch.org/short-reads/2023/07/24/8-facts-about-americans-with-disabilities/> (<https://www.pewresearch.org/short-reads/2023/07/24/8-facts-about-americans-with-disabilities/>)

Navigate to ▼

[51] Lever Team. 2023. *7 ChatGPT use cases for talent acquisition teams*. <https://www.lever.co/blog/chatgpt-use-cases/> (<https://www.lever.co/blog/chatgpt-use-cases/>)

Navigate to ▼

[52] Yannick l'Horty, Naomie Mahmoudi, Pascale Petit, and François-Charles Wolff. 2022. Is disability more discriminatory in hiring than ethnicity, address or gender? Evidence from a multi-criteria correspondence experiment. *Social Science & Medicine* 303 (2022), 114990.

Navigate to ▼

[53] Louis Lippens, Siel Vermeiren, and Stijn Baert. 2023. The state of hiring discrimination: A meta-analysis of (almost) all recent correspondence experiments. *European Economic Review* 151 (2023), 104315.

Navigate to ▼

[54] Brent J Lyons, Larry R Martinez, Enrica N Ruggs, Michelle R Hebl, Ann Marie Ryan, Katherine R O'Brien, and Adam Roebuck. 2018. To say or not to say: Different strategies of acknowledging a visible disability. *Journal of Management* 44, 5 (2018), 1980–2007.

Navigate to ▼

[55] Jennifer Elizabeth Marshall, Colm Fearon, Marianne Highwood, and Katy Warden. 2020. “What should I say to my employer...if anything?” – My disability disclosure dilemma. *International Journal of Educational Management* 34, 7 (2020), 1105–1117.

Navigate to ▼

[56] Bert Massie. 1994. Disabled people and social justice. Number 12. Institute for Public Policy Research.

Navigate to ▼

[57] Daniel McDuff, Shuang Ma, Yale Song, and Ashish Kapoor. 2019. Characterizing bias in classifiers using generative models. *Advances in neural information processing systems* 32 (2019).

Navigate to ▼

[58] Paige McGlaulin. 2023. Men like Elon Musk and Kanye West show why women and neurodiverse people struggle to break the leadership glass ceiling. *Fortune* (12 May 2023). <https://fortune.com/2023/05/12/men-elon-musk-kanye-west-women-neurodiversity-leadership-glass-ceiling/> (<https://fortune.com/2023/05/12/men-elon-musk-kanye-west-women-neurodiversity-leadership-glass-ceiling/>)

Navigate to ▼

[59] Katelyn Mei, Sonia Fereidooni, and Aylin Caliskan. 2023. Bias against 93 stigmatized groups in masked language models and downstream sentiment classification tasks. In *Proceedings of the 2023 ACM Conference on Fairness, Accountability, and Transparency*. 1699–1710.

Navigate to ▼

[60] Emma Mishel. 2016. Discrimination against queer women in the US workforce: A résumé audit study. *Socius* 2 (2016), 2378023115621316.

Navigate to ▼

[61] Haley Moss. 2020. Screened out onscreen: Disability discrimination, hiring bias, and artificial intelligence. *Denv. L. Rev.* 98 (2020), 775.

Navigate to ▼

[62] Dena F Mujtaba and Nihar R Mahapatra. 2019. Ethical considerations in AI-based recruitment. In *2019 IEEE International Symposium on Technology and Society (ISTAS)*. IEEE, 1–7.

Navigate to ▼

[63] Daniel Nemirovsky, Nicolas Thiebaut, Ye Xu, and Abhishek Gupta. 2021. Providing actionable feedback in hiring marketplaces using generative adversarial networks. In *Proceedings of the 14th ACM International Conference on Web Search and Data Mining*. 1089–1092.

Navigate to ▼

[64] Selin E Nugent and Susan Scott-Parker. 2022. Recruitment AI has a disability problem: Anticipating and mitigating unfair automated hiring decisions. In *Towards Trustworthy Artificial Intelligent Systems*. Springer, 85–96.

Navigate to ▼

[65] Jisoo Ock. 2022. The practical impact of bias against minority group applicants in resume screening on personnel selection outcomes. *Sustainability* 14, 15 (2022), 9438.

Navigate to ▼

[66] OpenAI. 2023. Introducing GPTs. *OpenAI Blog* (2023). <https://openai.com/blog/introducing-gpts/> (<https://openai.com/blog/introducing-gpts/>)

Navigate to ▼

[67] OpenAI. 2023. Introducing the GPT store. *OpenAI Blog* (2023). <https://openai.com/blog/introducing-the-gpt-store> (<https://openai.com/blog/introducing-the-gpt-store>)

Navigate to ▼

[68] Philip Oreopoulos and Diane Dechief. 2012. Why do some employers prefer to interview Matthew, but not Samir? New evidence from Toronto, Montreal, and Vancouver. *New Evidence from Toronto, Montreal, and Vancouver (February 2012)* (2012).

Navigate to ▼

[69] Kaja Larsen Østerud. 2023. Disability discrimination: Employer considerations of disabled jobseekers in light of the ideal worker. *Work, Employment and Society* 37, 3 (2023), 740–756.

Navigate to ▼

[70] European Parliament. 2023. EU AI act: First regulation on artificial intelligence. *European Parliament News* (2023). <https://www.europarl.europa.eu/news/en/headlines/society/20230601STO93804/eu-ai-act-first-regulation-on-artificial-intelligence> (<https://www.europarl.europa.eu/news/en/headlines/society/20230601STO93804/eu-ai-act-first-regulation-on-artificial-intelligence>) Updated on December 19, 2023.

Navigate to ▼

[71] Alejandro Pena, Ignacio Serna, Aythami Morales, and Julian Fierrez. 2020. Bias in multimodal AI: Testbed for fair automatic recruitment. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops*. 28–29.

Navigate to ▼

[72] Tammy Prater and Sara Bliss Kiser. 2002. Lies, lies, and more lies. *SAM Advanced Management Journal* 67, 2 (2002), 9.

Navigate to ▼

[73] Nitin Rane. 2023. Role and challenges of ChatGPT and similar generative artificial intelligence in human resource management. *Available at SSRN 4603230* (2023).

Navigate to ▼

[74] Chaminda Rathnayake and Aruni Gunawardana. 2023. The role of generative AI in enhancing human resource management recruitment, training, and performance evaluation perspectives. *International Journal of Social Analytics* 8, 11 (2023), 13–22.

Navigate to ▼

[75] Jean-François Ravaud, Béatrice Madiot, and Isabelle Ville. 1992. Discrimination towards disabled people seeking employment. *Social Science & Medicine* 35, 8 (1992), 951–958.

Navigate to ▼

[76] Alene Rhea, Kelsey Markey, Lauren D'Arinzo, Hilke Schellmann, Mona Sloane, Paul Squires, and Julia Stoyanovich. 2022. Resume format, LinkedIn URLs and other unexpected influences on AI personality prediction in hiring: Results of an audit. In *Proceedings of the 2022 AAAI/ACM Conference on AI, Ethics, and Society*. 572–587.

Navigate to ▼

[77] Emily Sheng, Josh Arnold, Zhou Yu, Kai-Wei Chang, and Nanyun Peng. 2021. Revealing persona biases in dialogue systems. *arXiv preprint arXiv:2104.08728* (<https://arXiv.org/abs/2104.08728>) (2021).

Navigate to ▼

[78] Emily Sheng, Kai-Wei Chang, Premkumar Natarajan, and Nanyun Peng. 2019. The woman worked as a babysitter: On biases in language generation. *arXiv preprint arXiv:1909.01326* (<https://arXiv.org/abs/1909.01326>) (2019).

Navigate to ▼

[79] Eric Michael Smith and Adina Williams. 2021. Hi, my name is Martha: Using names to measure and mitigate bias in generative dialogue models. *arXiv preprint arXiv:2109.03300* (<https://arXiv.org/abs/2109.03300>) (2021).

Navigate to ▼

[80] Calvin Solomon. 2020. Autism and employment: Implications for employers and adults with ASD. *Journal of Autism and Developmental Disorders* 50, 11 (2020), 4209–4217.

Navigate to ▼

[81] Alissa C. Stevens, Dianna D. Carroll, Elizabeth A. Courtney-Long, Qing C. Zhang, Michelle L. Sloan, Shannon Griffin-Blake, and Georgina Peacock. 2016. Adults with one or more functional disabilities—United States, 2011–2014. *Morbidity and Mortality Weekly Report* 65, 38 (2016), 1021–1025.

Navigate to ▼

[82] Luhang Sun, Mian Wei, Yibing Sun, Yoo Ji Suh, Liwei Shen, and Sijia Yang. 2023. Smiling women pitching down: Auditing representational and presentational gender biases in image generative AI. *arXiv preprint arXiv:2305.10566* (<https://arXiv.org/abs/2305.10566>) (2023).

Navigate to ▼

[83] Siri Thanasombat and John Trasviña. 2005. Screening names instead of qualifications: Testing with emailed resumes reveals racial preferences. *AAPI Nexus: Policy, Practice and Community* 3, 2 (2005), 105–

115.

Navigate to ▼

[84]

Nicholas Tilmes. 2022. Disability, fairness, and algorithmic bias in AI recruitment. *Ethics and Information Technology* 24, 2 (2022), 21.

Navigate to ▼

[85]

Thanh Tung Tran, Truong Giang Nguyen, Thai Hoa Dang, and Yuta Yoshinaga. 2023. Improving human resources’ efficiency with a generative AI-based resume analysis solution. In *International Conference on Future Data and Security Engineering*. Springer, 352–365.

Navigate to ▼

[86]

Swanand Vaishampayan, Sahar Farzanehpour, and Chris Brown. 2023. Procedural justice and fairness in automated resume parsers for tech hiring: Insights from candidate perspectives. In *2023 IEEE Symposium on Visual Languages and Human-Centric Computing (VL/HCC)*. IEEE, 103–108.

Navigate to ▼

[87]

Elmira Van den Broek, Anastasia Sergeeva, and Marleen Huysman. 2021. When the machine meets the expert: An ethnography of developing AI for hiring.*MIS quarterly* 45, 3 (2021).

Navigate to ▼

[88]

Akshaj Kumar Veldanda, Fabian Grob, Shailja Thakur, Hammond Pearce, Benjamin Tan, Ramesh Karri, and Siddharth Garg. 2023. Are Emily and Greg still more employable than Lakisha and Jamal? Investigating algorithmic hiring bias in the era of ChatGPT. *arXiv preprint arXiv:2310.05135*. (<https://arXiv.org/abs/2310.05135>) (2023).

Navigate to ▼

[89]

Pranav Narayanan Venkit, Mukund Srinath, and Shomir Wilson. 2022. A study of implicit bias in pretrained language models against people with disabilities. In *Proceedings of the 29th International Conference on Computational Linguistics*. 1324–1332.

Navigate to ▼

[90]

Neelie Verlinden. [n. d.]. ChatGPT for recruiting. <https://www.aihr.com/blog/chatgpt-for-recruiting/> (<https://www.aihr.com/blog/chatgpt-for-recruiting/>) Visited 5/2/2024.

Navigate to ▼

[91]

Bernadette Walker. 2024. OpenAI's GPT Shows Bias in Résumé Ranking Experiment. *Bloomberg* (8 March 2024). <https://www.bloomberg.com/news/newsletters/2024-03-08/openai-s-gpt-shows-bias-in-resume-screening-experiment-big-take> (<https://www.bloomberg.com/news/newsletters/2024-03-08/openai-s-gpt-shows-bias-in-resume-screening-experiment-big-take>) Accessed on 2024-05-01.

Navigate to ▼

[92]

Dana Wilson-Kovacs, Michelle K Ryan, S Alexander Haslam, and Anna Rabinovich. 2008. ‘Just because you can get a wheelchair in the building doesn't necessarily mean that you can still participate’: Barriers to the career advancement of disabled professionals. *Disability & Society* 23, 7 (2008), 705–717.

Navigate to ▼

[93]

Chantelle Wood and Megan Freeth. 2016. Students’ stereotypes of autism. *Journal of Educational Issues* 2, 2 (2016), 131–140.

Navigate to ▼

[94]

Michelle Yin, Dahlia Shaewitz, and Mahlet Megra. 2014. An uneven playing field: The lack of equal pay for people with disabilities. *Washington: American Institutes for Research* (2014), 1–17.

Navigate to ▼

[95]

Zapier. [n. d.]. Zapier integration: Generate candidate summaries in recruit CRM using ChatGPT. <https://zapier.com/apps/recruitcrm/integrations/recruitcrm/1199802/generate-summaries-via-chatgpt-for-new-talent-pool-candidates-in-recruit-crm> (<https://zapier.com/apps/recruitcrm/integrations/recruitcrm/1199802/generate-summaries-via-chatgpt-for-new-talent-pool-candidates-in-recruit-crm>) Integration page for Zapier and Recruit CRM, leveraging ChatGPT for automatic candidate summaries. Additional information may be available on the Recruit CRM and ChatGPT websites..

Navigate to ▼

[96]

A. Zarb. 2022. Assessing the role of gender in hiring: a field experiment on labour market discrimination. *SN Business & Economics* 2 (2022), 191.

Navigate to ▼

[97]

Mi Zhou, Vibhanshu Abhishek, and Kannan Srinivasan. 2023. Bias in generative AI (Work in Progress).

Navigate to ▼

A EXPERIMENT MATERIALS

This appendix includes an anonymized version of the job description and an anonymized representation of the resume used in the experiment.

A.1 Job Description

Minimum qualifications: Currently enrolled in a PhD degree in Computer Science, Linguistics, Statistics, Biostatistics, Applied Mathematics, Operations Research, Economics, Natural Sciences, or related technical field. Experience in one area of computer science (e.g., Natural Language Understanding, Computer Vision, Machine Learning, Deep Learning, Algorithmic Foundations of Optimization, Quantum Information Science, Data Science, Software Engineering, or similar areas). Preferred qualifications: Currently enrolled in a full-time degree program and returning to the program after completion of the internship. Currently attending a degree program in the US. Experience as a researcher, including internships, full-time, or at a lab. Experience contributing to research communities or efforts, including publishing papers in major conferences or journals. Experience with one or more general purpose programming languages (e.g., Python, Java, JavaScript, C/C++, etc.). Ability to communicate in English fluently. About the job The Student Researcher Program's primary objective is to foster academic collaborations with students through research at [COMPANY]. Join us for a paid Student Researcher position that offers the opportunity to work directly with [COMPANY] research scientists and engineers on research projects. The Student Researcher Program offers more opportunities for research students to work on critical research projects at [COMPANY] in a less structured way. The program allows opportunities beyond the limitations of our traditional internship program on aspects such as duration, time commitment, and working location (with options for on-site or remote). The topics student researchers work on tend to be open-ended and exploratory, and don't always have a clear deliverable like a traditional internship would. [COMPANY] Research is building the next generation of intelligent systems for all [COMPANY] products. To achieve this, we're working on projects that utilize the latest computer science techniques developed by skilled software engineers and research scientists. [COMPANY] Research teams collaborate closely with other teams across [COMPANY], maintaining the flexibility and versatility required to adapt new projects and foci that meet the demands of the world's fast-paced business needs. The US base salary range for this full-time position is 106,000 – 141,000. Our salary ranges are determined by role, level, and location. The range displayed on each job posting reflects the minimum and maximum target for new hire salaries for the position across all US locations. Within the range, individual pay is determined by work location and additional factors, including job-related skills, experience, and relevant education or training. Your recruiter can share more about the specific salary range for your preferred location during the hiring process. Please note that the compensation details listed in US role postings reflect the base salary only, and do not include bonus, equity, or benefits. Learn more about benefits at [COMPANY].

A.2 Jobseeker Resume Representation

The CV was seven pages long in PDF format, and ten pages long in text format. The resumes contained forty-nine resume items. The disability-enhanced resumes contained four extra, disability-related items. Below, we show layouts of the control resumes and enhanced resumes, highlighting the positions of the added items.




Figure 2: Control (CV) Resume Representation




Figure 3: Enhanced (ECV) Resume with Disability Representation

B POST-HOC TESTS

Post-hoc tests were conducted to address the following question: Will GPT rank resumes with additional awards that are not disability-related lower than those without?

B.1 Non-Disabled Award Tests

We ran post-hoc tests in Spring 2024 using GPT-4, replicating similar tests we performed in early Winter 2023 when deciding methodologies. We used non-disability dimensions, modifying the same four ECV resume items: [Var: No Dimension, Athlete, Seattle].

Table 5: Number of times the ECV was ranked first out of 10 trials with GPT-4. **Denotes statistical significance difference using Fisher's Exact test one-tailed test $p<0.05$*

Dimensions Tested	Number of Trials	ECV Ranked 1 st (GPT-4)
No Dimension Award	10	10*
Athletics Award	10	6
Regional Seattle Award	10	7

All of the award CVs ranked higher overall. Similar to our baseline experiment, we observed ‘ties’, which were absent in disability resume rankings. Unlike with the disability ECVs, GPT-4 acknowledged the extra items and ranked the new ECVs first based on them: “While still highly relevant and impressive, [CV] is essentially a subset of [ECV]...it lacks the additional details...” ... “[ECV] slightly edges out because it includes additional information in the ‘Awards and Honors’ section”. A limitation of this post-hoc test is that it used a different version of GPT-4, since GPT-4 had been updated after our data was collected in Winter 2023.

FOOTNOTE

*Both authors contributed equally to this research.



This work is licensed under a [Creative Commons Attribution International 4.0 License \(https://creativecommons.org/licenses/by/4.0/\)](https://creativecommons.org/licenses/by/4.0/).

FAccT '24, June 03–06, 2024, Rio de Janeiro, Brazil

© 2024 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-0450-5/24/06.
DOI: <https://doi.org/10.1145/3630106.3658933> (<https://doi.org/10.1145/3630106.3658933>)