

Name: Sudha Sree, Yerramsetty

Title of chapter: Chapter 09 - Distribution Shift

Summary:

This chapter delves into distribution shift, the phenomenon in which the data a model encounters upon deployment varies from the data on which it was trained, resulting in decreased performance. The chapter emphasizes that most machine learning algorithms assume that training and testing data are from the same distribution, but in practice, this assumption frequently fails owing to changes in environments, users, or collection techniques. It divides distribution shift into three types: covariate shift (when input features change), prior probability shift (when class/label proportions differ), and concept drift (when the relationship between inputs and outputs changes). Each type presents various obstacles and requires different detection and adaptation strategies.

Furthermore, the chapter explains how to detect, measure, and mitigate distribution shifts. It asserts that adaptation is possible even when we have a limited amount of deployment data. It discusses techniques such as domain adaptation, which recalibrates the model's decision threshold to adapt to prior probability shift, reweighting samples using importance weighting to adapt to covariate shift, and continual learning, which includes periodic retraining or online learning and model monitoring to adapt to concept drift. It also emphasizes how shifts can result from social or operational contexts (for example, models trained in one demographic group or geographical location may underperform in another). The chapter stresses that dealing with distribution shift is more than simply a technical challenge; it is also a concern of AI systems' resilience, fairness, and dependability. Understanding how and why data changes over time enables practitioners to create models that are reliable and successful in dynamic, real-world scenarios.

What are two points that you learned after reading this chapter?:

- Detecting and correcting distribution shifts, such as covariate shift, prior probability shift, and concept drift, requires different detection and correction strategies. Recognizing which type is occurring is critical for resolving performance degradation.
- While distribution shift is inevitable in real-world scenarios, it may be handled by continual monitoring, retraining, and adaptive learning to preserve model reliability over time.

What are two questions that you have about this chapter?:

- How can we design automated systems that continuously detect and adapt to distribution shifts without relying heavily on labeled data?
- What are the most effective strategies to distinguish between harmful distribution shifts and natural variations that do not significantly impact model performance?

Do you have some initial thoughts or insights to approach the questions?:

- To design automated adaptive systems, semi-supervised learning or online learning methods could be leveraged to detect shifts through changes in feature distributions or prediction confidence without requiring many new labels. Additionally, drift detection algorithms like DDM or ADWIN can trigger updates when a performance drop is detected.
- Setting monitoring thresholds and evaluating model uncertainty can assist to distinguish between harmful and harmless shifts. Using performance metrics on holdout validation sets from recent data can demonstrate whether a shift has a meaningful impact on outcomes, signalling whether model retraining is required.