

Name: Sudha Sree, Yerramsetty

Title of chapter: Chapter 12 - Interpretability and Explainability

Summary:

This chapter explores how to make machine learning models understandable and trustworthy for different types of users. It emphasizes that interpretability focuses on understanding how a model works internally, while explainability focuses on conveying these insights to people in useful ways. The chapter uses the Hilo platform, which predicts housing prices and creditworthiness, to demonstrate why transparency is essential in real-world applications involving people's lives and finances. It discusses several explanation methods, distinguishing between local (local explanations interpret specific predictions) vs global (global explanations describe the model's behavior overall), approximate (exact explanations show the true decision logic, which is possible only in simple models) vs exact (approximate explanations create interpretable surrogates for complex models), and feature-based (feature-based methods focus on the importance of individual features) vs sample-based approaches (sample-based approaches use representative samples to explain model behavior). Each method is suited to a specific "persona": model builders who need to debug and refine models, decision-makers who must justify decisions, and affected users who seek actionable recourse when decisions impact them directly.

The chapter also addresses specific strategies, such as Partial Dependence Plots (PDP) (PDP shows the average effect of a feature across the dataset) and Individual Conditional Expectation (ICE) (ICE shows how a feature affects individual predictions) for model builders, LIME (Local Interpretable Model-agnostic Explanations) (LIME approximates a complex model locally around one specific prediction using a simple linear model, showing which features contributed most to that prediction) for decision-makers and the Contrastive Explanations Method (CEM) (CEM is related to adversarial examples but is designed to produce interpretable and sparse feature adjustments rather than imperceptible perturbations; it produces both local and exact explanations) for end users, demonstrating how these approaches may help make opaque models easier to understand. Finally, it explores quantifying interpretability using three levels of evaluation: application-grounded, human-grounded, and functionally-grounded, each with differing degrees of realism and cost. The main takeaway is that true explainability necessitates not only technological tools but also human-centered evaluation. Therefore, the chapter warns that functionally-grounded evaluations alone are insufficient. Before deploying models, teams should conduct at least some human or application-grounded studies to ensure explanations are genuinely useful to people. Effective interpretability should help build trust, accountability, and fairness while ensuring that explanations meet the actual requirements of those affected by AI decisions.

What are two points that you learned after reading this chapter?:

- Different explanation methods are suited for different audiences. Developers, decision-makers, and end users each require different explanations tailored to their roles and goals.

- Interpretability evaluation is not just quantitative; human-centered evaluations (such as application-grounded studies) are crucial for ensuring that the explanations remain pertinent and understandable.

What are two questions that you have about this chapter?:

- How can we balance the trade-off between model performance and interpretability when high accuracy often comes from complex, less transparent models?
- Can contrastive explanations such as CEM be scaled consistently to accommodate big, real-world datasets with hundreds of characteristics and dynamic user inputs?

Do you have some initial thoughts or insights to approach the questions?:

- A potential method to solve this is exploring hybrid modeling strategies, which involve employing interpretable models when practical and explainability tools (such as LIME or SHAP) to clarify complicated models when required. In this way, we preserve accuracy while maintaining transparency.
- Future research could focus on computing efficiency and feature selection to simplify the contrastive explanation process. Combining CEM with dimensionality reduction or model distillation might make it more suitable for real-world usage.