

Name: Sudha Sree, Yerramsetty

Title of paper: “Big Data’s Disparate Impact” by Solon Barocas & Andrew D. Selbst

Summary:

This paper examines how big data and algorithmic decision-making can unintentionally reproduce and reinforce patterns of social inequality and discrimination. While data mining is often portrayed as an objective, bias-free method, the authors demonstrate that algorithms are only as fair as the data and design choices that underlies them. Because data often reflects historical human biases or institutional discrimination, algorithms trained on it can yield results that disproportionately hurt marginalized groups, even if there is no intention to discriminate. These discriminatory results become more difficult to identify, explain, or challenge because they occur as statistical byproducts rather than deliberate human decisions.

Using Title VII of the Civil Rights Act as a legal lens, the authors argue that US anti-discrimination law is inadequate to address this new type of algorithmic bias. Under existing law, many data-driven techniques are legally justified as "business necessities" provided they properly predict outcomes such as job performance, even if these predictions reflect historical biases. The paper concludes that solving algorithmic discrimination would require a reassessment of the legal and moral notions of fairness and discrimination, rather than just removing evident biases. The authors highlight both the potential and risk of big data, underlining the need for further technological, legal, and political changes to ensure fairness in automated decision-making.

Two points that you like about the paper (strengths):

- The research paper effectively connects the technical and legal perspectives, demonstrating how computer science techniques (such as data collecting and model training) intersect with anti-discrimination law. This interdisciplinary strategy strengthens and contextualizes the argument for policymakers and technologists alike.
- The authors go beyond the typical emphasis on human bias to illustrate how discrimination might develop naturally from data and algorithms. This is a sophisticated and innovative concept that helps to rethink fairness in an automated society.

Two points that you did not like about the paper (weaknesses):

- Despite a detailed analysis of the problem, the research paper provides only a few specific policy frameworks or strategies to reduce algorithmic discrimination in real-world systems. The issue is deeply understood by readers such as myself, but we are left with no idea of how to effectively handle it.
- The legal analysis is limited to the context of discrimination during the hiring process in the workplace. With examples from a wider range of business sectors or legal systems, the authors' ideas may have been more broadly relevant, even if they touch on issues for housing, credit, and other areas only in passing.

Detailed comments and questions:

- The paper's distinction between intentional and unintentional discrimination is insightful, but it raises a question: if bias is embedded in data, who should be held accountable: the data collectors, model developers, or the institutions deploying the algorithms?
- Although the authors discuss how variables can act as proxies for protected attributes, is it possible that technological methods like algorithmic auditing or fairness-aware machine learning could offer at least some partial solutions to this problem?
- A profound conflict in the way society defines fairness is brought to light by the debate between anticlassification and antistatutory concepts. Should anti-discrimination laws continue to focus on process-based fairness (ensuring neutral procedures) or evolve toward outcome-based fairness (directly addressing inequality)?
- I appreciated the authors' acknowledgment that data mining can also reduce discrimination when done responsibly. It would be interesting to explore examples or frameworks where algorithms have demonstrably improved fairness in hiring, lending, or criminal justice.