

Name: Sudha Sree, Yerramsetty

Title of paper: “Representation Bias in Data: A Survey on Identification and Resolution Techniques” by Nima Shahbazi et al.

Summary:

This paper provides an extensive overview of representation bias, which describes how imbalanced or incomplete data can lead to unfair results in machine learning systems. It argues that representation bias occurs when specific groups or features are underrepresented or misrepresented in datasets, resulting in models that exacerbate and worsen existing socioeconomic disparities. The authors structure the survey around two key goals: identifying and resolving representation bias. They explore these tasks across both structured data (such as tabular or relational databases) and unstructured data (text, images, speech, and graphs). They outline key methods, fairness measures, and tactics for each type of data, including mainly reweighting, data augmentation, adversarial training, and embedding debiasing.

The study additionally proposes taxonomies for categorizing fairness strategies based on data modality and methodological approach, which will help readers understand how solutions in one area (e.g., image recognition) might inspire approaches in another. The authors emphasize the increasing importance of investigating bias in graphs, which provide unique issues due to their non-independent, non-Euclidean structure. The study concludes by describing future research topics, such as the need for fairness methodologies that can be applied to new data types (such as streaming or spatiotemporal data) and the creation of improved metrics to measure representation bias.

Two points that you like about the paper (strengths):

- I think that the paper provides a clear and organized taxonomy that categorizes bias identification and mitigation methods by data type and level (graph-level, embedding-level, etc.), which makes this complex field more understandable and navigable for the readers.
- It covers a wide range of modalities, starting from structured tabular data to unstructured data like images, text, speech, and graphs, offering a complete overview of how representation bias manifests differently across data domains.

Two points that you did not like about the paper (weaknesses):

- While the paper identifies promising future research directions, it provides limited discussion on real-world implementation challenges or how these techniques can be integrated into production-level AI systems.
- This study is mostly a survey of existing literature so it tends to skim over each literature instead of explaining any of the algorithms in depth. I couldn't fully understand a few of the algorithms being used.

Detailed comments and questions:

- The paper's mathematical formulation of graph-based bias indicators (e.g., assortative mixing coefficient and representation bias) is useful, but readers may benefit from more intuitive explanations or a few examples to help them understand what the metrics imply in reality.
- It would be interesting to know if combining various resolution measures (e.g., preprocessing and embedding-level debiasing) results in better fairness outcomes than employing one strategy alone.
- Overall, I found the study to be a very useful resource for understanding the landscape of fairness research and I understood that it laid the groundwork for more consistent and practical debiasing approaches. The taxonomy diagrams really helped me get a visual picture of how all of these methodologies can be classified in a hierarchical manner.