

Name: Sudha Sree, Yerramsetty

Title of paper: “Fairness in Ranking, Part I: Score-Based Ranking” by Meike Zehlike et al.

Summary:

This paper provides a comprehensive survey of the research landscape on fairness in ranking, focusing specifically on score-based ranking, where items (usually people) are ranked based on numeric scores or formulas. Unlike classification, ranking is a relative decision mechanism that focuses on moving one candidate up or down which affects everyone else, which in turn introduces unique fairness challenges. The authors unify concepts across multiple research communities (information retrieval, databases, algorithms, recommender systems) and establish a structured way to think about fairness in ranking. They introduce four classification frameworks: value-based ethical foundations, bias typology, fairness mitigation objectives, and intervention types (pre-/in-/post-processing). The core of the survey explains how fairness constraints such as proportional representation, prefix fairness, and exposure fairness can be applied to score-based ranking systems.

The paper also shows how data bias, historical disadvantage, and measurement issues can result in unfair rankings and how algorithms can mitigate these issues while balancing utility and fairness. It uses a rich running example of university admissions to illustrate how different fairness criteria impact real-world ranking outcomes. The authors review technical methods such as score adjustments, fairness-aware top-k selection, fairness-constrained re-ranking, and exposure-based fairness techniques from information retrieval. They conclude Part I by consolidating insights and preparing readers for Part II, which covers fairness in supervised learning-to-rank and recommender/matching systems.

Two points that you like about the paper (strengths):

- The paper systematically organizes a highly fragmented body of fairness-in-ranking research into coherent frameworks. This makes it much easier to understand relationships between fairness goals, ethical motivations, and algorithmic interventions.
- The admissions example effectively illustrates the complexity of fairness constraints (e.g., top-k parity vs. prefix parity). It grounds abstract fairness ideas in a scenario that readers can intuitively understand.

Two points that you did not like about the paper (weaknesses):

- Although the survey covers many algorithmic interventions, it often describes them conceptually rather than with full mathematical rigor. Readers seeking precise formulation may need to consult the original papers.
- The paper acknowledges causal fairness but focuses primarily on associational fairness. Given the importance of causal reasoning for understanding structural bias, a deeper integration of causal frameworks would have strengthened the discussion and piqued my interest more.

Detailed comments and questions:

- Value frameworks are extremely helpful, but their mapping to specific algorithms could have been more explicit. For example, when a fairness objective is tied to an ethical principle (like anti-discrimination), the survey could explicitly list which algorithmic methods best embody that principle.
- The distinction between top-k fairness and prefix fairness is excellent but raises practical questions. Do institutions actually need prefix fairness in practice, or is top-k parity sufficient? The paper could expand on real-world applicability.
- The dataset discussion is useful but somewhat limited. More commentary on limitations of existing datasets, for example, demographic sparsity or availability constraints would help practitioners choose datasets wisely.
- Future work could include standardized benchmarks for fair ranking. Since classification fairness has mature benchmarks, ranking fairness could benefit from a similar effort.